

CSC411: Project #2

Due on Sunday, February 18, 2018

Hao Zhang & Tyler Gamvrelis

February 5, 2018

Foreword

In this project, neural networks of various depth were used to recognize handwritten digits and faces.

Part 1 provides a description of the MNIST dataset from which the images of handwritten digits were taken. Part 2 implements the computation of a simple network. Part 3 presents the *sum of the negative log-probabilities* as a cost function, and derives the expression for its gradient with respect to one of its weights. Part 4 details the training and optimization procedures for a digit-recognition neural network, and part 5 improves upon the results by modifying gradient descent to use momentum. Learning curves are presents in parts 4 and 5. Part 6 presents an analysis of network behaviour with respect to two of its weights. Part 7 provides an analysis of the performance of two different backpropagation computation techniques. Part 8 presents a face recognition network architecture for classifying actors, and uses PyTorch for implementation. Part 9 presents visualizations of hidden unit weights relevant to two of the actors. Part 10 uses activations of AlexNet to train a neural network to perform classification of the actors.

System Details for Reproducibility:

- Python 2.7.14
- Libraries:
 - numpy
 - matplotlib
 - pylab
 - time
 - os
 - scipy
 - urllib
 - cPickle
 - PyTorch

Part 1

Dataset description

The MNIST dataset is made of thousands of 28 by 28 pixel images of the handwritten digits: 0 to 9. The images are split into training set and test set images labelled ‘train0’ to ‘train9’ and ‘test0’ to ‘test9’. The number of images with each label is as follows:

| Label | Number of Images | Label | Number of Images |
|--------|------------------|-------|------------------|
| train0 | 5923 | test0 | 980 |
| train1 | 6742 | test0 | 1135 |
| train2 | 5958 | test0 | 1032 |
| train3 | 6131 | test0 | 1010 |
| train4 | 5842 | test0 | 982 |
| train5 | 5421 | test0 | 892 |
| train6 | 5918 | test0 | 958 |
| train7 | 6265 | test0 | 1028 |
| train8 | 5851 | test0 | 974 |
| train9 | 5949 | test0 | 1009 |

Ten images of each number were taken from the training sets and displayed in Figure 1. The correct labels of most of the pictures can be discerned at a glance by humans. However, since the digits are handwritten, some of them may not be completely obvious. For example, Figure 1cv is categorized as a 9 but looks like an 8.



Figure 1: Subset of the MNIST dataset.

Part 2

Computing a simple network

In this part, the simple network depicted in Figure 2 was implemented as a function in Python using NumPy, the code listing for which is presented in Figure 3.

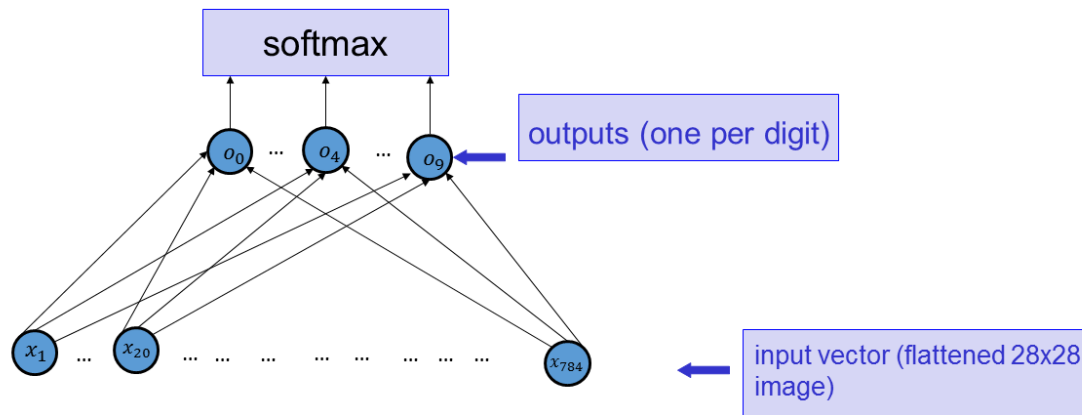


Figure 2: Simple network diagram from project handout.

```

1 def softmax(y):
2     '''
3     Return the output of the softmax function for the matrix of output y. y
4     is an NxM matrix where N is the number of outputs for a single case, and M
5     is the number of cases
6     '''
7     return exp(y)/tile(sum(exp(y),0), (len(y),1))
8
9 def Part2(theta, X):
10    '''
11    Part2 returns the vectorized multiplication of the (n x 10) parameter matrix
12    theta with the data X.
13
14    Arguments:
15    theta -- (n x 10) matrix of parameters (weights and biases)
16    x -- (n x 1) matrix whose rows correspond to pixels in images
17    '''
18
19    return softmax(np.dot(theta.T, X))

```

Figure 3: Python implementation of network using NumPy.

Part 3

Part 4

Part 5

Part 6

Part 7

Part 8

Part 9

Part 10