CSC411: Project #2

Due on Sunday, February 18, 2018

Hao Zhang & Tyler Gamvrelis

February 9, 2018

Foreword

In this project, neural networks of various depth were used to recognize handwritten digits and faces.

Part 1 provides a description of the MNIST dataset from which the images of handwritten digits were taken. Part 2 implements the computation of a simple network. Part 3 presents the *sum of the negative log-probabilities* as a cost function, and derives the expression for its gradient with respect to one of its weights. Part 4 details the training and optimization procedures for a digit-recognition neural network, and part 5 improves upon the results by modifying gradient descent to use momentum. Learning curves are presents in parts 4 and 5. Part 6 presents an analysis of network behaviour with respect to two of its weights. Part 7 provides an analysis of the performance of two different backpropagation computation techniques. Part 8 presents a face recognition network architecture for classifying actors, and uses PyTorch for implementation. Part 9 presents visualizations of hidden unit weights relevant to two of the actors. Part 10 uses activations of AlexNet to train a neural network to perform classification of the actors.

System Details for Reproducibility:

- Python 2.7.14

- Libraries:

    - numpy
    - matplotlib
    - pylab
    - time
    - os
    - scipy
    - urllib
    - cPickle
    - PyTorch

Part 1

Dataset description

The MNIST dataset is made of thousands of 28 by 28 pixel images of the handwritten digits: 0 to 9. The images are split into training set and test set images labelled 'train0' to 'train9' and 'test0' to 'test9'. The number of images with each label is presented in Table 1, below.

Label	Number of Images	Label	Number of Images
train0	5923	test0	980
train1	6742	test1	1135
train2	5958	test2	1032
train3	6131	test3	1010
train4	5842	test4	982
train5	5421	test5	892
train6	5918	test6	958
train7	6265	test7	1028
train8	5851	test8	974
train9	5949	test9	1009

Table 1: Quantity of each type of image in the MNIST dataset.

Ten images of each number were taken from the training sets and displayed in Figure 1. The correct labels of most of the pictures can be discerned at a glance by humans However, since the digits are handwritten, some of them may not be completely obvious. For example, Figure 1cv is categorized as a 9 but looks like an 8.

Figure 1: Subset of the MNIST dataset.

Part 2

Computing a Simple Network

In this part, the simple network depicted in Figure 2 was implemented as a function in Python using NumPy, the code listing for which is presented is Figure 3.
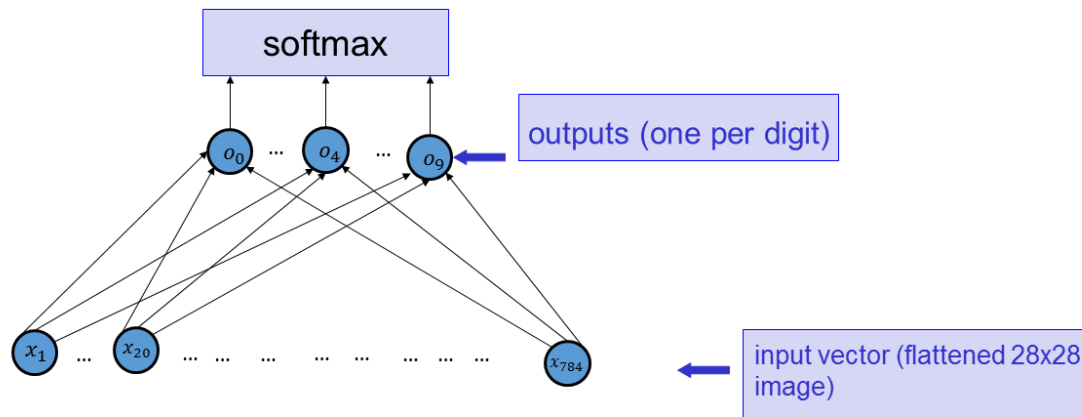

Figure 2: Simple network diagram from project handout.

```python
def softmax(y):
    '''
    Return the output of the softmax function for the matrix of output y. y
    is an NxM matrix where N is the number of outputs for a single case, and M
    is the number of cases
    '''
    return exp(y)/tile(sum(exp(y),0), (len(y),1))

def SimpleNetwork(W, X):
    '''
    SimpleNetwork returns the vectorized multiplication of the (n x 10)
    parameter matrix W with the data X.
    Arguments:
        W -- (n x 10) matrix of parameters (weights and biases)
        x -- (n x m) matrix whose i-th column corresponds to the i-th training
    image
    '''
    return softmax(np.dot(W.T, X))
```

Figure 3: Python implementation of network using NumPy.

Part 3

Cost Function of Negative Log Probabilities
The Cost function that will be used for the network described in part 2 is:

$$C = -\sum_{q=1}^{m}(\sum_{l=1}^{k} y_l log(p_l))_q$$

Where k is the number of output nodes, m is the number of training samples, $p_l = \frac{e^{o_l}}{\sum_{c=1}^{k} e^{o_c}}$ and y_l is equal to 1 if the training sample is labelled as l and 0 otherwise.

Partial Derivative of the Cost Function (3a)

Let the matrix $W = \begin{bmatrix} w_{11} & w_{12} & w_{13} & \dots & w_{1k} \\ w_{21} & w_{22} & w_{23} & \dots & w_{2k} \\ \dots\dots\dots\dots\dots\dots\dots\dots \\ w_{n1} & w_{n2} & w_{n3} & \dots & w_{nk} \end{bmatrix}$ and let W_j be the j^{th} column of matrix W. This results

in $o_j = W_j^T x$ where $x = \begin{bmatrix} x_1 \\ x_2 \\ .. \\ x_n \end{bmatrix}$ and n is equal to the number of pixels with an extra 1 to multiply into the

bias. That is, x_1^q is always taken to be 1 while $x_2^q, ..., x_n^q$ are the inputs to the network corresponding to the q^{th} training example.

To calculate the partial derivative of C, we will consider only **a single training sample**: $C = -\sum_{l=1}^{k} y_l log(p_l)$
By the chain rule, the partial derivative with respect to w_{ij} is:

$$\frac{\partial C(W)}{\partial w_{ij}} = -\sum_{l=1}^{k} \frac{\partial C}{\partial p_l} \frac{\partial p_l}{\partial o_j} \frac{\partial o_j}{\partial w_{ij}}$$

The first term is straight forward: $\frac{\partial C}{\partial p_l} = \frac{y_l}{p_l}$.

For the second term, we have $\frac{\partial p_l}{\partial o_j} = \begin{cases} \frac{-e^{o_l}e^{o_j}}{(\sum_{c=1}^{k} e^{o_c})^2} = -p_l p_j, \text{if } l \neq j \\ \frac{e^{o_j}}{\sum_{c=1}^{k} e^{o_c}} - \frac{e^{2o_j}}{(\sum_{c=1}^{k} e^{o_c})^2} = p_j - p_j^2, \text{if } l = j \end{cases}$

From the equation: $o_j = W_j^T x$, we have $o_j = w_{1j}x_1 + w_{2j}x_2 + \dots + w_{nj}x_n$ and it is clear to see that: $\frac{\partial o_j}{\partial w_{ij}} = x_i$.

Putting it all together, we have:

$$\frac{\partial C(W)}{\partial w_{ij}} = -(\sum_{l \neq j}(-\frac{y_l}{p_l}p_l p_j) + \frac{y_j}{p_j}(p_j - p_j^2))x_i = (\sum_{l \neq j}(y_l p_j) + y_j(p_j - 1))x_i$$

Considering that we are using 1-hot encoding, only a single y_l can equal 1. Therefore,

$$\frac{\partial C(W)}{\partial w_{ij}} = \begin{cases} (p_j - 1)x_i, & \text{if } y_j = 1 \\ p_j x_i, & \text{otherwise} \end{cases}$$

Which is equivalent to saying: $\frac{\partial C(W)}{\partial w_{ij}} = (p_j - y_j)x_i$. To extend this to all training samples, we just have to sum up each individual contribution, resulting in the equation: $\frac{\partial C(W)}{\partial w_{ij}} = \sum_{q=1}^{m}((p_j - y_j)x_i)_q$.

Vectorized Implementation of Gradient (3b)

To vectorize the gradient, we first define the gradient matrix to be:

$$\frac{\partial C(W)}{\partial W} = \begin{bmatrix} \frac{\partial C}{w_{11}} & \cdots & \cdots & \cdots & \frac{\partial C}{w_{1k}} \\ \frac{\partial C}{w_{21}} & \cdots & \cdots & \cdots & \frac{\partial C}{w_{2k}} \\ \dots\dots\dots\dots\dots\dots\dots\dots \\ \frac{\partial C}{w_{n1}} & \cdots & \cdots & \cdots & \frac{\partial C}{w_{nk}} \end{bmatrix} = \begin{bmatrix} \sum_{q=1}^{m}((p_1-y_1)x_1)_q & \cdots & \cdots & \cdots & \sum_{q=1}^{m}((p_k-y_k)x_1)_q \\ \sum_{q=1}^{m}((p_1-y_1)x_2)_q & \cdots & \cdots & \cdots & \sum_{q=1}^{m}((p_k-y_k)x_2)_q \\ \dots\dots\dots\dots\dots\dots\dots\dots\dots\dots\dots\dots\dots\dots\dots\dots\dots \\ \sum_{q=1}^{m}((p_1-y_1)x_n)_q & \cdots & \cdots & \cdots & \sum_{q=1}^{m}((p_k-y_k)x_n)_q \end{bmatrix}$$

This is equivalent to: $X(P-Y)^T$ where:

$$X = \begin{bmatrix} x_1^{(1)} & x_1^{(2)} & x_1^{(3)} & \cdots & x_1^{(m)} \\ x_2^{(1)} & x_2^{(2)} & x_2^{(3)} & \cdots & x_2^{(m)} \\ \dots\dots\dots\dots\dots\dots\dots\dots \\ x_n^{(1)} & x_n^{(2)} & x_n^{(3)} & \cdots & x_n^{(m)} \end{bmatrix} \quad Y = \begin{bmatrix} y_1^{(1)} & y_1^{(2)} & y_1^{(3)} & \cdots & y_1^{(m)} \\ y_2^{(1)} & y_2^{(2)} & y_2^{(3)} & \cdots & y_2^{(m)} \\ \dots\dots\dots\dots\dots\dots\dots\dots \\ y_k^{(1)} & y_k^{(2)} & y_k^{(3)} & \cdots & y_k^{(m)} \end{bmatrix} \quad P = \begin{bmatrix} p_1^{(1)} & p_1^{(2)} & p_1^{(3)} & \cdots & p_1^{(m)} \\ p_2^{(1)} & p_2^{(2)} & p_2^{(3)} & \cdots & p_2^{(m)} \\ \dots\dots\dots\dots\dots\dots\dots\dots \\ p_k^{(1)} & p_k^{(2)} & p_k^{(3)} & \cdots & p_k^{(m)} \end{bmatrix}$$

Where the superscript of each element represents the index of the training samples. The code used to compute the gradient is shown in Figure 4.

```python
def negLogLossGrad(X, Y, W):
    '''
    negLogLossGrad returns the gradient of the cost function of negative log
    probabilities.

    Arguments:
        X -- Input image(s) from which predictions will be made (n x m)
        Y -- Target output(s) (actual labels) (10 x m)
        W -- Weight matrix (n x 10)
    '''

    P = SimpleNetwork(W, X) # Get the prediction for X using the weights W

    return np.dot(X, (P - Y).transpose())
```

Figure 4: Vectorized Implementation of the Negative Log Loss Gradient

Finite Difference Approximation (3b)

To check that the gradient was computed correctly, the gradient with respect to weight w_{ij} was approximated using finite differences, for all coordinates in the weight matrix. 10 values for the differential quantity, h, were tested, for which the maximum total error was 2.698 for $h = 1$ and the minimum total error was 3.054×10^{-7} for $h = 1 \times 10^{-7}$. Thus, for sufficiently small h, the finite difference approximation and vectorized computation converge to the same quantity, suggesting the vectorized computation is correct. The test case used to obtain this result is shown in Figure 5. The complete results are summarized in Table 2, below.

Total error	h
2.698202291518717	1.0
0.200166359876926	0.1
0.019912135073950066	0.01
0.0019906216532900034	0.001
0.00019905646908402463	0.0001
1.990372311189148e-05	1e-05
1.9841905873896337e-06	1e-06
3.2450363651737035e-07	1e-07
3.0542104362263345e-06	1e-08
2.747131061797692e-05	1e-09

Table 2: Sum of differences between vectorized gradient computation and finite difference approximation for 10 h-values.

```
# Test part 3
X = [[1, 2, 3, 4], [1, 2, 3, 4], [1, 2, 3, 4], [1, 2, 3 ,4],[1, 2, 3, 4]]
X = np.array(X)
W = [[1, 0, 1.2, 3], [1, 2, 0.2, 1.2], [1, 8, 1, 4], [1, -2, 3, 1], [1, 0, 3, 1]]
W = np.array(W)
Y = [[1, 0, 0, 0], [0, 1, 0, 0], [0, 0, 1, 0], [0, 0, 0, 1]]
Y = np.array(Y)
G1 = negLogLossGrad(X, Y, W)
h = 1.0
for i in range(10):
    G2 = negLogLossGrad_FiniteDifference(X, Y, W, h)
    print "Total error: ", sum(abs(G1 - G2)), "h: ", h
    h /= 10
```

Figure 5: Test Case Used to Calculate Error Between Finite Difference Approximation and the Vectorized Gradient

Part 4

In this section, the neural network from part 2 was trained using vanilla gradient descent. The optimization procedure is described below, and learning curves and visualizations of the weights going into each of the output units are presented.

Optimization Procedure
Words

Learning Curves
Words

Weights Going Into Output Units
Words

Part 5

Part 6

Part 7

Part 8

Part 9

Part 10